

Alzheimers-Detection

August 31, 2023

Alzheimers detection dataset obtained from <https://www.kaggle.com/datasets/taeefnajib/handwriting-data-to-detect-alzheimers-disease>.

```
[ ]: # Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: # Loading data
alz = pd.read_csv("alzheimers.csv")
```

1 Exploratory Data Analysis

```
[ ]: # First 5 rows of data
alz.head(5)
```

```
[ ]:
```

	ID	air_time1	disp_index1	gmrt_in_air1	gmrt_on_paper1	\
0	id_1	5160	0.000013	120.804174	86.853334	
1	id_2	51980	0.000016	115.318238	83.448681	
2	id_3	2600	0.000010	229.933997	172.761858	
3	id_4	2130	0.000010	369.403342	183.193104	
4	id_5	2310	0.000007	257.997131	111.275889	

	max_x_extension1	max_y_extension1	mean_acc_in_air1	mean_acc_on_paper1	\
0	957	6601	0.361800	0.217459	
1	1694	6998	0.272513	0.144880	
2	2333	5802	0.387020	0.181342	
3	1756	8159	0.556879	0.164502	
4	987	4732	0.266077	0.145104	

	mean_gmrt1	...	mean_jerk_in_air25	mean_jerk_on_paper25	\
0	103.828754	...	0.141434	0.024471	
1	99.383459	...	0.049663	0.018368	
2	201.347928	...	0.178194	0.017174	
3	276.298223	...	0.113905	0.019860	
4	184.636510	...	0.121782	0.020872	

	mean_speed_in_air25	mean_speed_on_paper25	num_of_pendown25	paper_time25	\
0	5.596487	3.184589	71	40120	
1	1.665973	0.950249	129	126700	
2	4.000781	2.392521	74	45480	
3	4.206746	1.613522	123	67945	
4	3.319036	1.680629	92	37285	

	pressure_mean25	pressure_var25	total_time25	class
0	1749.278166	296102.7676	144605	P
1	1504.768272	278744.2850	298640	P
2	1431.443492	144411.7055	79025	P
3	1465.843329	230184.7154	181220	P
4	1841.702561	158290.0255	72575	P

[5 rows x 452 columns]

```
[ ]: # Shape of data
alz.shape
```

```
[ ]: (174, 452)
```

```
[ ]: # Data information
alz.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 174 entries, 0 to 173
Columns: 452 entries, ID to class
dtypes: float64(300), int64(150), object(2)
memory usage: 614.6+ KB
```

```
[ ]: # Checking for object column names
alz.select_dtypes(include = "object").columns.tolist()
```

```
[ ]: ['ID', 'class']
```

```
[ ]: # Checking for missing values
alz.isna().sum() # No NA values
```

```
[ ]: ID          0
     air_time1   0
     disp_index1 0
     gmrt_in_air1 0
     gmrt_on_paper1 0
     ..
     paper_time25 0
     pressure_mean25 0
     pressure_var25 0
```

```
total_time25      0
class             0
Length: 452, dtype: int64
```

2 Feature Engineering

```
[ ]: # Removing ID column
alz = alz.drop("ID", axis = 1)
alz.head(5)
```

```
[ ]:   air_time1  disp_index1  gmrt_in_air1  gmrt_on_paper1  max_x_extension1  \
0         5160      0.000013    120.804174      86.853334             957
1        51980      0.000016    115.318238      83.448681            1694
2         2600      0.000010    229.933997     172.761858            2333
3         2130      0.000010    369.403342     183.193104            1756
4         2310      0.000007    257.997131     111.275889             987

      max_y_extension1  mean_acc_in_air1  mean_acc_on_paper1  mean_gmrt1  \
0                6601          0.361800          0.217459  103.828754
1                6998          0.272513          0.144880   99.383459
2                5802          0.387020          0.181342  201.347928
3                8159          0.556879          0.164502  276.298223
4                4732          0.266077          0.145104  184.636510

      mean_jerk_in_air1  ...  mean_jerk_in_air25  mean_jerk_on_paper25  \
0          0.051836  ...              0.141434              0.024471
1          0.039827  ...              0.049663              0.018368
2          0.064220  ...              0.178194              0.017174
3          0.090408  ...              0.113905              0.019860
4          0.037528  ...              0.121782              0.020872

      mean_speed_in_air25  mean_speed_on_paper25  num_of_pendown25  paper_time25  \
0           5.596487              3.184589              71          40120
1           1.665973              0.950249              129          126700
2           4.000781              2.392521              74          45480
3           4.206746              1.613522              123          67945
4           3.319036              1.680629              92          37285

      pressure_mean25  pressure_var25  total_time25  class
0          1749.278166      296102.7676          144605      P
1          1504.768272      278744.2850          298640      P
2          1431.443492      144411.7055           79025      P
3          1465.843329      230184.7154          181220      P
4          1841.702561      158290.0255           72575      P
```

```
[5 rows x 451 columns]
```

```
[ ]: # Converting class to numeric
alz["class"] = alz["class"].replace({'P': 1, 'H': 0})
alz["class"]
```

```
[ ]: 0      1
      1      1
      2      1
      3      1
      4      1
      ..
     169     0
     170     0
     171     0
     172     0
     173     0
Name: class, Length: 174, dtype: int64
```

3 Model Training

```
[ ]: from sklearn.model_selection import train_test_split

# Separating features from target
X = alz.drop(columns=["class"])
y = alz["class"]

# Training data with a 70/30 split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.7,
↳random_state = 42)
```

```
[ ]: # Random Forest
from sklearn.ensemble import RandomForestClassifier
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.tree import plot_tree

# Creating random forest pipeline with scaled data
pipe = Pipeline([
    ('scaler', StandardScaler()),
    ('classifier', RandomForestClassifier(random_state = 42, max_samples = 0.6,
↳min_samples_leaf = 2))
])

# Fitting pipeline
pipe.fit(X_train, y_train)

# Predicting target values
```

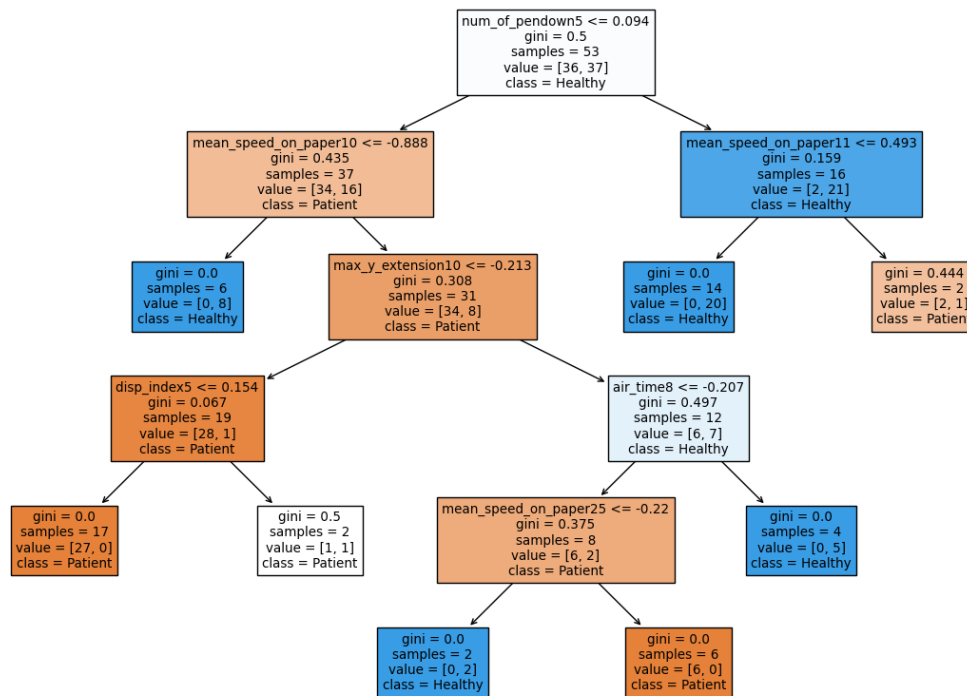
```
y_pred = pipe.predict(X_test)
```

```
[ ]: # Plotting first tree in the random forest
tree_viz = pipe.named_steps['classifier'].estimators_[0]

fig, ax = plt.subplots(figsize = (15, 10))

plot_tree(tree_viz, feature_names = alz.columns.tolist(), class_names = ["Patient", "Healthy"], filled = True)

plt.show()
```

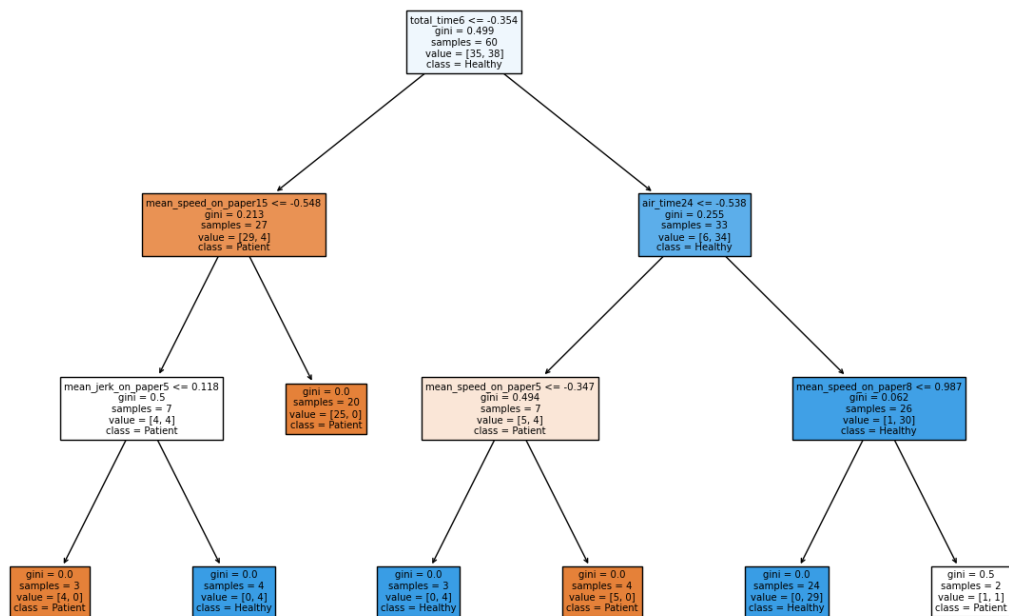


```
[ ]: # Plotting fiftieth tree in the random forest
tree_viz = pipe.named_steps['classifier'].estimators_[49]

fig, ax = plt.subplots(figsize = (15, 10))

plot_tree(tree_viz, feature_names = alz.columns.tolist(), class_names = ["Patient", "Healthy"], filled = True)

plt.show()
```



4 Results

```
[ ]: from sklearn.metrics import f1_score

# F1 score is high so this random forest model is a good predictor of the target
f1 = f1_score(y_test, y_pred)
print("F1 Score:", f1)
```

F1 Score: 0.9019607843137256

```
[ ]: from sklearn.metrics import roc_auc_score, roc_curve

# False positive and true positive rates
fpr, tpr, _ = roc_curve(y_test, y_pred)

# AUC
auc = roc_auc_score(y_test, y_pred)

# Plotting ROC curve
fig, ax = plt.subplots()

ax.plot(fpr, tpr, color = 'darkorange', lw = 2, label = 'ROC curve (AUC = {:.2f})'.format(auc))
```

```


ax.plot([0, 1], [0, 1], color = 'navy', lw = 2, linestyle = '--')

ax.set_xlim([0.0, 1.0])
ax.set_ylim([0.0, 1.05])
ax.set_xlabel('False Positive Rate')
ax.set_ylabel('True Positive Rate')
ax.set_title('Receiver Operating Characteristic (ROC) Curve')
ax.legend(loc = "lower right")

sns.despine()

plt.show()

```



```

[ ]: from sklearn.metrics import confusion_matrix

# Creating confusion matrix
conf_matrix = confusion_matrix(y_test,y_pred)

# Plotting confusion matrix
fig, ax = plt.subplots()

```


```

sns.heatmap(conf_matrix,
            annot = True,
            fmt = 'g',
            xticklabels = ['Positive', 'Negative'],
            yticklabels = ['Positive', 'Negative'],
            cmap = ["Red", "Green", "Red", "Green"],
            cbar = False,
            annot_kws = {"size": 20},
            ax = ax)

ax.set_title('Confusion Matrix', fontsize = 17)
ax.set_ylabel('Prediction', fontsize = 13)
ax.set_xlabel('Actual', fontsize = 13)

plt.show()

```



```

[ ]: # Creating TP/FP/TN/FN
TP = conf_matrix[1, 1]

```



```
FN = conf_matrix[1, 0]
TN = conf_matrix[0, 0]
FP = conf_matrix[0, 1]

# Printing results of predictions
accuracy = (TP + TN) / (TP + TN + FP + FN)
precision = (TP) / (TP + FP)
sensitivity = TP / (TP + FN)
specificity = TN / (TN + FP)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Sensitivity:", sensitivity)
print("Specificity:", specificity)
```

```
Accuracy: 0.9056603773584906
Precision: 0.8846153846153846
Sensitivity: 0.92
Specificity: 0.8928571428571429
```